

Sep 08, 22 1:10	main.py	Page 1/4
<pre>#!/bin/env python3.8 """ Homework Assignment #1: Gregory Presser Help Recived From: Husam Almanakly """ import os import logging import matplotlib import matplotlib.pyplot as plt import numpy as np import tensorflow as tf from absl import app from absl import flags from tqdm import trange from dataclasses import dataclass, field, InitVar script_path = os.path.dirname(os.path.realpath(__file__)) LOWER_VAL = 0 UPPER_VAL = 1 @dataclass class LinearModel: weights: np.ndarray bias: float mew: np.ndarray sigma: np.ndarray @dataclass class Data: model: LinearModel rng: InitVar[np.random.Generator] num_features: int num_samples: int sigma: float x: np.ndarray = field(init=False) y: np.ndarray = field(init=False) def __post_init__(self, rng): self.index = np.arange(self.num_samples) self.x = rng.uniform(LOWER_VAL, UPPER_VAL, size=(self.num_samples, 1)) clean_y = np.sin(2 * np.pi * self.x) self.y = clean_y + rng.normal(loc=0, scale=0.1, size=(self.num_samples, 1)) def get_batch(self, rng, batch_size): """ Select random subset of examples for training batch """ choices = rng.choice(self.index, size=batch_size) return self.x[choices], self.y[choices].flatten() def compare_linear_models(a: LinearModel, b: LinearModel): for w_a, w_b in zip(a.weights, b.weights): print(f"{w_a:0.2f}, {w_b:0.2f}")</pre>		

Sep 08, 22 1:10	main.py	Page 2/4
<pre>print(f"{a.bias:0.2f}, {b.bias:0.2f}") font = { # "family": "Adobe Caslon Pro", "size": 10, } matplotlib.style.use("classic") matplotlib.rc("font", **font) FLAGS = flags.FLAGS flags.DEFINE_integer("num_features", 4, "Number of features in record") flags.DEFINE_integer("num_samples", 50, "Number of samples in dataset") flags.DEFINE_integer("batch_size", 16, "Number of samples in batch") flags.DEFINE_integer("num_iters", 300, "Number of SGD iterations") flags.DEFINE_float("learning_rate", 0.01, "Learning rate / step size for SGD") flags.DEFINE_integer("random_seed", 31415, "Random seed") flags.DEFINE_float("sigma_noise", 0.1, "Standard deviation of noise random variable") flags.DEFINE_bool("debug", True, "Set logging level to debug") class Model(tf.Module): def __init__(self, rng, num_features): """ A plain linear regression model with a bias term """ self.num_features = num_features self.b = tf.Variable(tf.zeros(shape=[1, 1]), name="bias") self.w = tf.Variable(rng.normal(shape=[self.num_features, 1]), name="weights") self.mew = tf.Variable(tf.cast(tf.linspace(LOWER_VAL, UPPER_VAL, self.num_features), tf.float32), name="mew",) self.sigma = (tf.Variable(tf.ones(shape=[self.num_features, 1]),) * 0.3) def __call__(self, x): gaussians = tf.transpose(self.w) * tf.math.exp(-((x - tf.transpose(self.mew)) ** 2 / (tf.transpose(self.sigma) ** 2))) return tf.squeeze(tf.reduce_sum(gaussians, 1) + self.b) @property def model(self): return LinearModel(self.w.numpy().reshape([self.num_features]), self.b.numpy().squeeze(), self.mew.numpy().reshape([self.num_features]), self.sigma.numpy().reshape([self.num_features]),) def main(a):</pre>		

Sep 08, 22 1:10

main.py

Page 3/4

```

logging.basicConfig()

if FLAGS.debug:
    logging.getLogger().setLevel(logging.DEBUG)

# Safe np and tf PRNG
seed_sequence = np.random.SeedSequence(FLAGS.random_seed)
np_seed, tf_seed = seed_sequence.spawn(2)
np_rng = np.random.default_rng(np_seed)
tf_rng = tf.random.Generator.from_seed(tf_seed.entropy)

data_generating_model = LinearModel(
    weights=np_rng.integers(low=0, high=5, size=(FLAGS.num_features)),
    bias=2,
    mew=np_rng.integers(low=0, high=1, size=(FLAGS.num_features)),
    sigma=np_rng.integers(low=0, high=1, size=(FLAGS.num_features)),
)
logging.debug(data_generating_model)

data = Data(
    data_generating_model,
    np_rng,
    FLAGS.num_features,
    FLAGS.num_samples,
    FLAGS.sigma_noise,
)

model = Model(tf_rng, FLAGS.num_features)
logging.debug(model.model)

optimizer = tf.optimizers.SGD(learning_rate=FLAGS.learning_rate)

bar = trange(FLAGS.num_iters)
for i in bar:
    with tf.GradientTape() as tape:
        x, y = data.get_batch(np_rng, FLAGS.batch_size)
        y_hat = model(x)
        loss = 0.5 * tf.reduce_mean((y_hat - y) ** 2)

        grads = tape.gradient(loss, model.trainable_variables)
        optimizer.apply_gradients(zip(grads, model.trainable_variables))

    bar.set_description(f"Loss @ {i} => {loss.numpy():0.6f}")
    bar.refresh()

logging.debug(model.model)

# print out true values versus estimates
print("w, w_hat")
compare_linear_models(data.model, model.model)

fig, ax = plt.subplots(1, 2, figsize=(10, 3), dpi=200)

ax[0].set_title("Sinewave Regression")
ax[0].set_xlabel("x")
ylab = ax[0].set_ylabel("y", labelpad=10)
ylab.set_rotation(0)

xs = np.linspace(np.amin(model.mew) * 1.5, np.amax(model.mew) * 1.5, 1000)
xs = xs[: np.newaxis]
y_hat = model(xs.reshape(1000, 1))
ax[0].plot(xs, np.squeeze(y_hat), "--", color="red")

```

Sep 08, 22 1:10

main.py

Page 4/4

```

ax[0].plot(np.squeeze(data.x), data.y, "o", color="blue")

real_y = np.sin(2 * np.pi * xs)
ax[0].plot(xs, real_y, color="green")

ax[1].set_title("Basis Functions")
ax[1].set_xlabel("x")
ylab = ax[1].set_ylabel("y", labelpad=10)
ylab.set_rotation(0)

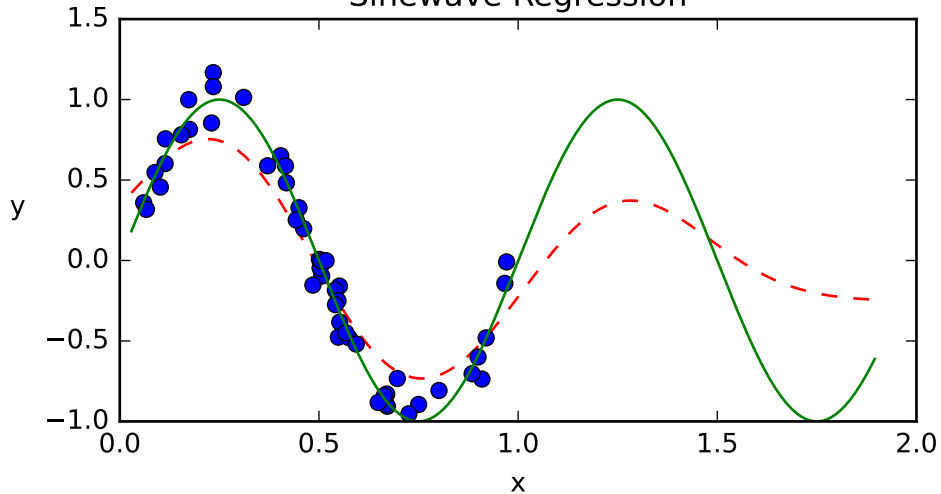
phi = np.zeros((1000, model.num_features))
for i in range(model.num_features):
    phi[:, i] = np.exp(-(xs.T - model.mew[i]) ** 2) / (model.sigma[i] ** 2)
)
ax[1].plot(xs, phi)

plt.tight_layout()
plt.savefig(f"{script_path}/fit.pdf")

if __name__ == "__main__":
    app.run(main)

```

Sinewave Regression



Basis Functions

